

Lecture 11 - Intro to Neural Networks

- I Logistic Regression
- II Neural Networks

Deep Learning

- Computational power
- Data available
- Algorithms

I Logistic Regression

goal 1.0 find cats in images { 1 → presence of cat
0 → absence of cat

$$64 \times 64 \text{ image} = \begin{pmatrix} k_1 \\ \vdots \\ k_n \end{pmatrix} \rightarrow w \cdot k + b$$

$$\hat{y} = G(w^T k) = G(w \cdot k + b)$$

$$(4, 128 \times 4) \rightarrow (64 \times 4 \times 3, 1)$$

i) initialize w, b

ii) find the optimal w, b

iii) Use $\hat{y} = G(w \cdot k + b)$ to predict

$$\mathcal{L} = -[y \log \hat{y} + (1-y) \log(1-\hat{y})]$$

$$\begin{cases} w = w - \alpha \frac{\partial \mathcal{L}}{\partial w} \\ b = b - \alpha \frac{\partial \mathcal{L}}{\partial b} \end{cases}$$

(Eq1) neuron: linear + activation

$$n = 64 \times 64 \times 3 + 1 \text{ params}$$

goal 2.0 find cat | iguana | lion in images

$$x = \begin{pmatrix} k_1 \\ \vdots \\ k_n \end{pmatrix} \rightarrow \hat{y}_1 = a_1^{[G1]} = G(w_1 \cdot k + b_1^{[G1]})$$

$$\rightarrow \hat{y}_2 = a_2^{[G1]} = G(w_2 \cdot k + b_2^{[G1]})$$

$$\rightarrow \hat{y}_3 = a_3^{[G1]} = G(w_3 \cdot k + b_3^{[G1]})$$

3n params

$$\boxed{y} \Rightarrow \begin{pmatrix} 1 \\ 1 \\ 0 \end{pmatrix} \begin{matrix} p(\text{cat}) \\ p(\text{lion}) \\ p(\text{lizard}) \end{matrix} \Rightarrow \text{robust, neurons don't talk to each other}$$

$$L_{ce} = - \sum_{k=1}^3 [y_k \log \hat{y}_k + (1-y_k) \log (1-\hat{y}_k)]$$

goal 3.0 + constraint

unique animal on an image

$$\boxed{y} = \begin{pmatrix} y_1 \\ \vdots \\ y_K \end{pmatrix} \quad \begin{array}{c} z_1^{(0)} \\ z_2^{(0)} \\ z_3^{(0)} \end{array} = e^{z_1^{(0)}} / \sum_{k=1}^3 e^{z_k^{(0)}} \\ = \dots \\ = \dots$$

→ softmax

multi-class
network

3 w paras $\binom{1}{0} \binom{0}{1} \binom{0}{1}$

Cross-entropy loss

$$L_{ce} = - \sum_{k=1}^3 y_k \log \hat{y}_k$$

? age of cat ↗ linear regression

ReLU - rectified linear units



II Neural Networks

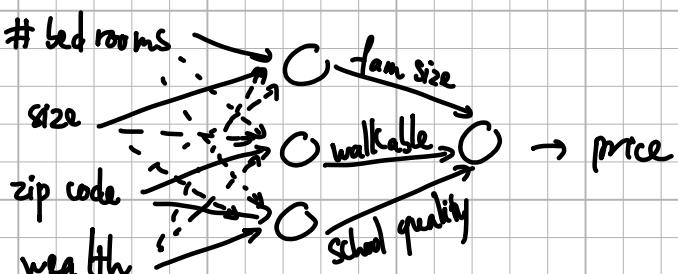
goal image \Rightarrow Cat vs no Cat

- neurons are not connected to each other
- cluster of neurons



→ fully connected layer \rightarrow paras $3n+3$ $\downarrow L \cdot 3 + 2$ $\downarrow 1 \cdot 2 + 1$ output layer = last
hidden layer

house price prediction



end to end learning | blackbox model

Propagation equations

$$(1,1) \rightarrow z^{[1]} = w^{[1]} x + b^{[1]} \quad \leftarrow (1,1)$$

$$(1,1) \rightarrow z^{[2]} = \sigma(z^{[1]}) \quad \leftarrow (1,1)$$

$$z^{[2]} = w^{[2]} a^{[1]} + b^{[2]}$$

$$a^{[2]} = \sigma(z^{[2]})$$

$$z^{[3]} = w^{[3]} a^{[2]} + b^{[3]}$$

$$\hat{y} = a^{[3]} = \sigma(z^{[3]})$$

• What happens for an input batch
of m examples?

$$X = \begin{pmatrix} | & | & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & | \end{pmatrix}$$

broadcasting
(numpy)

$$\begin{aligned} z^{[1]} &= w^{[1]} X + b^{[1]} \quad \leftarrow \begin{pmatrix} | & | & | \\ z^{(1)} & z^{(2)} & \dots & z^{(m)} \\ | & | & | \end{pmatrix} \quad \leftarrow (1,1) \\ z^{[m]} &= \begin{pmatrix} | & | & | \\ z^{(1)} & z^{(2)} & \dots & z^{(m)} \\ | & | & | \end{pmatrix} \quad \leftarrow \begin{pmatrix} | & | & | \\ b^{[1]} & b^{[2]} & \dots & b^{[m]} \\ | & | & | \end{pmatrix} \quad \leftarrow m \end{aligned}$$

optimizing $w^{[1]} \dots w^{[l]}, b^{[1]} \dots b^{[l]}$

Define loss / cost function

$$J(\hat{y}, y) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}^{(i)}$$

$$\text{with } \mathcal{L}^{(i)} = -[y^{(i)} \log \hat{y}^{(i)} + (1-y^{(i)}) \log (1-\hat{y}^{(i)})]$$

backward propagation

$$\forall l = 1 \dots 3 \quad \begin{cases} w^{[l]} = w^{[l]} - \alpha \frac{\partial J}{\partial w^{[l]}} \\ b^{[l]} = b^{[l]} - \alpha \frac{\partial J}{\partial b^{[l]}} \end{cases}$$

→ Chain rule

$$\frac{\partial J}{\partial w^{[3]}} = \underbrace{\frac{\partial J}{\partial z^{[3]}}} \cdot \underbrace{\frac{\partial a^{[3]}}{\partial z^{[3]}}} \cdot \underbrace{\frac{\partial z^{[3]}}{\partial w^{[3]}}}$$

$$\frac{\partial J}{\partial w^{[2]}} = \underbrace{\frac{\partial J}{\partial z^{[2]}}} \cdot \underbrace{\frac{\partial z^{[2]}}{\partial a^{[2]}}} \cdot \underbrace{\frac{\partial a^{[2]}}{\partial z^{[2]}}} \cdot \underbrace{\frac{\partial z^{[2]}}{\partial w^{[2]}}}$$

$$\frac{\partial J}{\partial w^{[1]}} = \underbrace{\frac{\partial J}{\partial z^{[1]}}} \cdot \underbrace{\frac{\partial z^{[1]}}{\partial a^{[1]}}} \cdot \underbrace{\frac{\partial a^{[1]}}{\partial z^{[1]}}} \cdot \underbrace{\frac{\partial z^{[1]}}{\partial w^{[1]}}}$$

Lecture 12 - Backprop & Improving Neural Networks

I log reg with a NN mindset

II Neural Networks

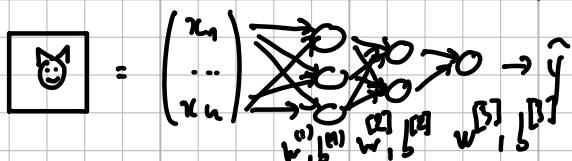
→ Back Propagation

III Improving your NNs

Backpropagation

$$\text{Cost function: } J(\hat{y}, y) = \frac{1}{m} \sum_{i=1}^m L^{(i)}(\hat{y}^{(i)}, y^{(i)})$$

$$\text{with } L^{(i)} = -[y^{(i)} \log \hat{y}^{(i)} + (1-y^{(i)}) \log (1-\hat{y}^{(i)})]$$



$$\text{Update } w^{[2]} = w^{[2]} - \alpha \frac{\partial J}{\partial w^{[2]}}$$

$$\frac{\partial L}{\partial w^{[2]}} = (a^{[1]} - y^{(i)}) a^{[2]T} = - (y^{(i)} - a^{[2]}) \underbrace{a^{[2]T}}_{\frac{\partial z^{[2]}}{\partial w^{[2]}}} \frac{\partial z^{[2]}}{\partial w^{[2]}}$$

$$\frac{\partial J}{\partial w^{[2]}} = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} - a^{[2]}) a^{[2]T}$$

Recall

$$g(x) = 1 - g(-x)$$

$$g(x) = \frac{1}{1+e^{-x}} = \frac{e^x}{e^x + 1}$$

$$g'(x) = g(x)(1-g(x))$$

$$\begin{aligned} \frac{\partial L}{\partial w^{[2]}} &= (a^{[2]} - y) \cdot w^{[2]T} a^{[2]} (1-a^{[2]}) a^{[1]T} \\ &= \cancel{w^{[2]T}} \times \cancel{a^{[2]}} \underbrace{(1-a^{[2]})}_{(1,2)^T} \underbrace{(a^{[1]} - y) \cdot a^{[1]T}}_{2,1} \end{aligned}$$

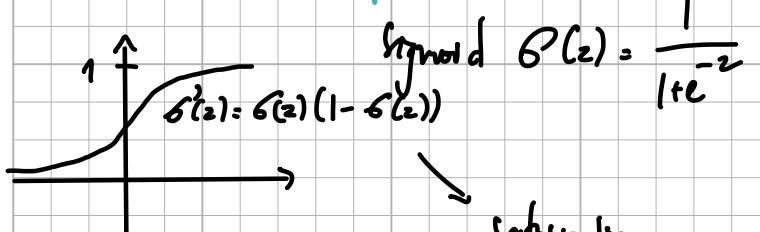
$$\hookrightarrow \frac{\partial J}{\partial w^{[2]}} = \frac{1}{m} \sum_{i=1}^m \frac{\partial L}{\partial w^{[2]}}$$

Cache → Computation efficiency

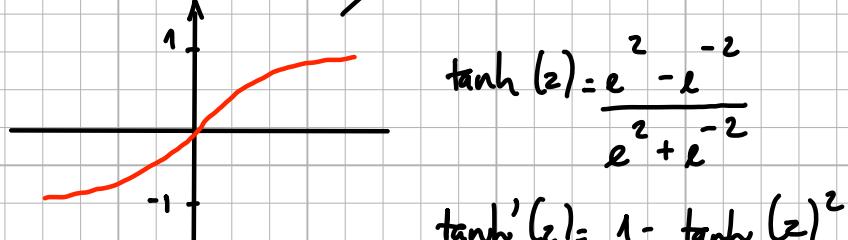
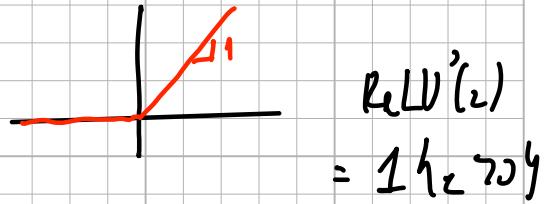
↳ Memory cost

III Improving your NNs

A Activation functions



$$\text{ReLU}(z) = \begin{cases} 0 & \text{if } z \leq 0 \\ 1 & \text{if } z > 0 \end{cases}$$



why do we need activation functions?

C

O

O

O

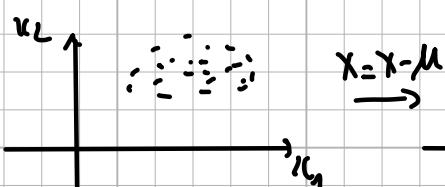
activations: $f_C: z \rightarrow z$

$$\begin{aligned} \hat{y} &= a^{[3]} = \frac{1}{2} = w^{[3]}_1 z^{[2]}_1 + b^{[3]} \\ &= w^{[3]}_1 w^{[2]}_1 (w^{[1]}_1 x + b^{[1]}) + w^{[3]}_2 z^{[2]}_2 + b^{[3]} \\ &= Wx + b \end{aligned}$$

B Initialization methods

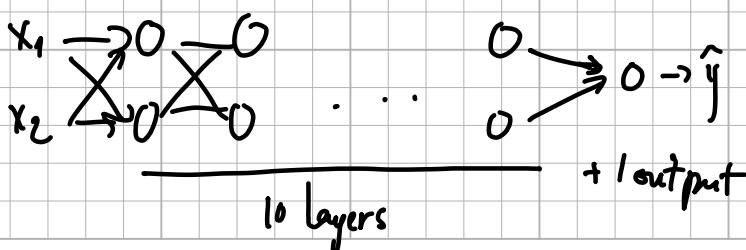
Normalizing your input

$$x = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$$

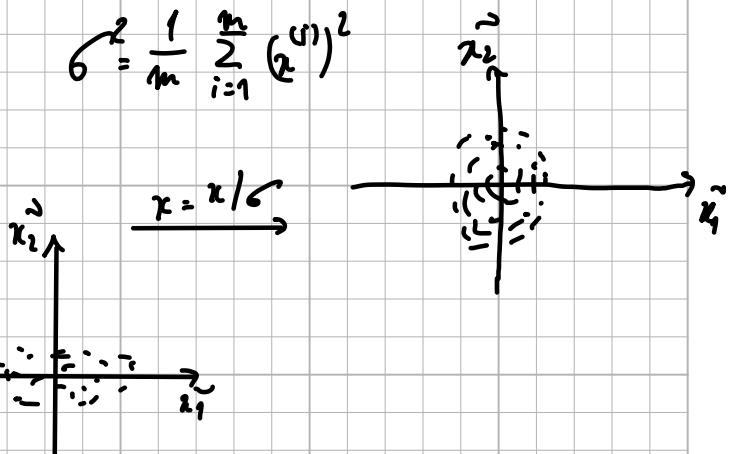


$$\begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix} = \mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

Vanishing | Exploding gradients



$$\sigma^2 = \frac{1}{m} \sum_{i=1}^m (x^{(i)})^2$$



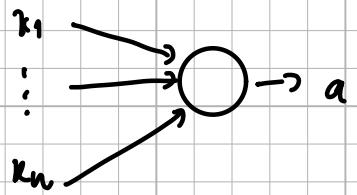
$$g: z \rightarrow z$$

$$L=0$$

$$\begin{aligned} \hat{y} &= w^{[L]} \cdot v^{[L-1]} \dots w^{[1]} \cdot x \\ &= w^{[L]} \cdot v^{[L-1]} \dots w^{[1]} \cdot x \end{aligned}$$

Example with 1 neuron :

$$a = \sigma(z)$$



$$z = w_1 z_1 + \dots + w_n z_n$$

large $n \rightarrow$ small w

$$w^{[l]} \sim \frac{1}{n} \quad w = \text{np.random. random (shape)} * \text{np.sqrt}\left(\frac{1}{n^{[L-1]}}\right)$$

Sigmoid , ReLU $\frac{c}{w^{[l-1]}}$

Xavier Initialization

$$w^{[l]} \sim \sqrt{\frac{1}{m^{[l-1]}}} \quad \text{for tanh}$$

He Initialization:

$$w^{[l]} \sim \sqrt{\frac{2}{m^{[l]} + n^{[l]}}}$$

C optimization

$$X = (x^{(1)} \ x^{(2)} \ \dots \ x^{(m)})$$

$$Y = (y^{(1)} \ y^{(2)} \ \dots \ y^{(m)})$$

$$\text{Batch } X = (X^{(1)} \ \dots \ X^{(100)})$$

$$(x^{(1)} \dots x^{(100)})$$

$$\text{Batch } Y = (Y^{(1)} \ \dots \ Y^{(100)})$$

$$J = \frac{1}{100} \sum_{i=1}^{100} L^{(i)}$$

Algo

for iteration $t = 1 \dots$

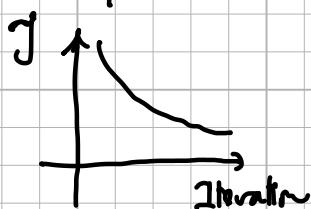
Select batch $(X^{(t)}, Y^{(t)})$

forward Prop \rightarrow Cost

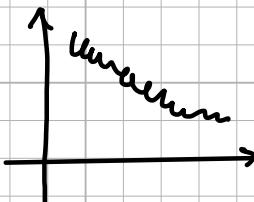
backprop Batch

Update $w^{[l]}, b^{[l]}$

Graph

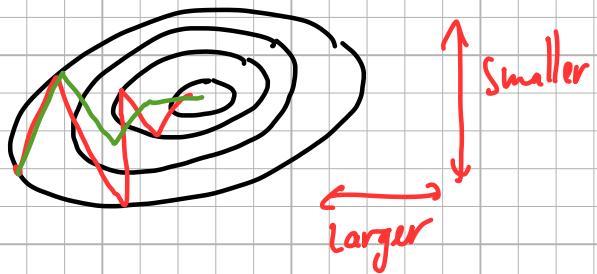


Batch Gradient
descent



Mini-batch
gradient descent

GD + Momentum Agle

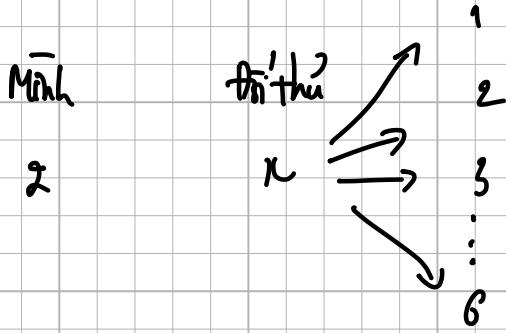


weights
↑

$$\left. \begin{array}{l} v = \beta v + (1-\beta) \frac{\partial L}{\partial w} \\ w = w - \alpha v \end{array} \right\}$$

Lecture 13 - Debugging ML Models and Error Analysis

claim có 1 cái ít nhất 2



→ Minh thừa [khi $x > 2$ và chỉ thứ đơn']

↳ 3, 4, 5, 6

→ Minh thừa [khi $x = 2$ và chỉ thứ đơn' 2 2s]

Các thi đơn' có thể đơn'

$$1 \rightarrow x \cdot 4(3 \rightarrow 6) = 4 \rightarrow 10 + h$$

$$2 \rightarrow x \cdot 6(1 \rightarrow 6) = 6$$

→ thi đơn' 2 [1.6] \ {2} → Minh reject → Minh thg
↳ Minh không reject



Minh claim $x_1 \geq y_1$

thi' claim $x_2 \geq y_2 \geq x_1$

→ optimal để đơn'

• 4 cái ít nhất y_1

→ thi' $x_2 \geq y_2 \geq x_1 = y_1$

$x_2 \geq y_2 \geq x_1 \geq y_1$

↓

→ thi' $x_2 \geq y_2 \geq x_1 \geq y_1$

moves = { $x \dots y$ }

$f[y_1]: y_1 >_r 3:$

Optimal to make this

$f[i][y]$

Minh (x_1) 1

1 2 3 4 5 6

Minh (x_2) 2

Minh claim: $x_1 > y_1$

Opponent's claim $x_2 > y_2 > y_1$

$\begin{cases} f(y_1) = \text{true} \text{ new (g) th}\bar{y} \\ f(y_2) = \text{false} \text{ v\bar{a} } y_2 > y_1 \\ f(y_1) = \begin{cases} \text{true, } \exists y_2 > y_1 : f[y_2] = \text{false} \\ \text{false, other} \end{cases} \end{cases}$

$f[i][y_1] = \begin{cases} \text{true, } \exists y_2 > y_1 : f[i][y_2] = \text{false} \\ \text{false, other} \end{cases}$

iswinning(state) {

moves = [1, 2, 3, 4, 5, 6]

for (all move in moves) {

newState = move if > state else return false

if (!iswinning(newState)) return True;

}

return false

}

→

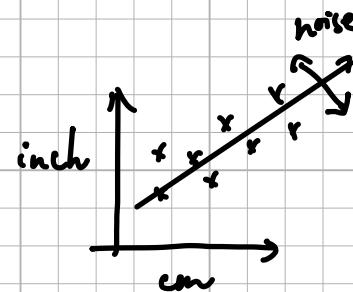
PCA - Principal Component Analysis → Unsupervised Learning

ICA - Independent Component Analysis

- Unlabeled dataset

$$\{x^{(1)}, \dots, x^{(n)}\} \in \mathbb{R}^n$$

Reduce dimension from $n \rightarrow k$
 $k < n$



Pre-processing:

$$\begin{aligned} 1. \mu &= \frac{1}{m} \sum_{i=1}^m x^{(i)} && \left. \begin{array}{l} \text{zero of} \\ \text{mean} \end{array} \right\} \\ 2. x^{(i)} &\leftarrow x^{(i)} - \mu && \\ 3. S_j^2 &:= \frac{1}{m} \sum_i x_j^{(i)T} x_j^{(i)} && \left. \begin{array}{l} \text{standardize} \\ \text{variance to 1} \end{array} \right\} \\ 4. x_j^{(i)} &\leftarrow \frac{x_j^{(i)}}{\sqrt{S_j^2}} \end{aligned}$$

$$\Rightarrow \text{Max } u^T \Sigma u$$

$u: \|u\|_2 = 1 \rightarrow u^T u = 1$

⇒ Sol: u is principal eigenvector of Σ

11 ways to derive PCA!

General Case: If you wish to project data to k -dim

→ Set u_1, \dots, u_k : top k -eigenvectors

Have $x^{(i)} \in \mathbb{R}^n$ (say $n = 1000$)

u_1, \dots, u_k (say $k = 60$)

$$x^{(i)} \rightarrow (u_1^T x^{(i)}, \dots, u_k^T x^{(i)}) = y^{(i)} \in \mathbb{R}^k$$

$$x^{(i)} \approx y_1^{(i)} u_1 + \dots + y_k^{(i)} u_k$$

- PCA → find direction of green line

$$\rightarrow \vec{u}, \|u\|_2 = 1$$

$$\begin{matrix} u^{(i)} \\ u^T u \end{matrix}$$

Choose u to max:

$$\text{Max } \frac{1}{m} \sum_{i=1}^m (x^{(i)T} u)^2$$

$$u: \|u\|_2 = 1 \quad = \frac{1}{m} \sum_{i=1}^m u^T x^{(i)} (x^{(i)T} u)$$

$$= u^T \underbrace{\left(\frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T} \right)}_{\Sigma} u$$

$$\Sigma = \frac{1}{m} \sum_{i=1}^m x^{(i)} x^{(i)T}$$

- Application:

- Visualization

Project from $n=7 \rightarrow 2, 3 = 2$

- Compression for ML efficiency

- Reduce overfitting

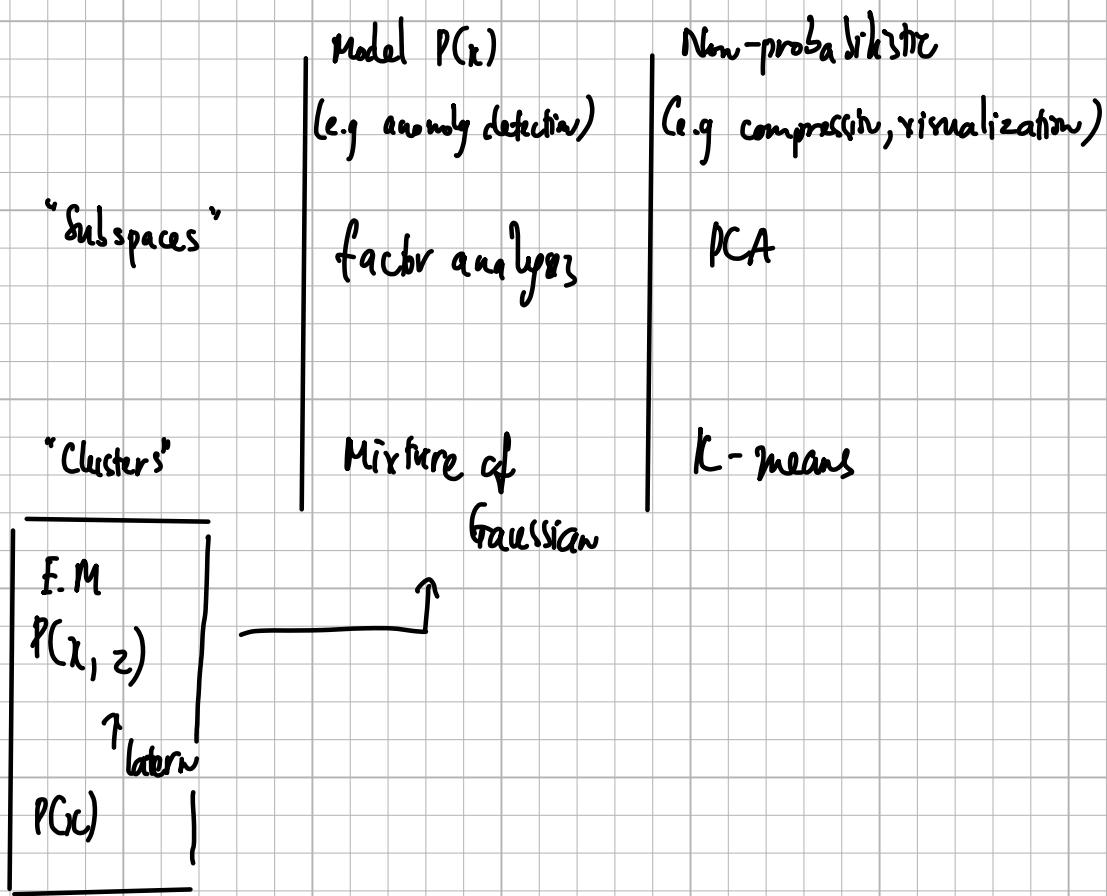
- outlier detection } matching

- Eigenfaces

Questionnaire

- Rules of thumb

- Before using PCA, use original data.



ICA

\cap
Independent

- Original sources: $s \in \mathbb{R}^n$ (n speakers)

$s_j^{(i)}$: signal from speaker j at time i

→ We observe

$$x^{(i)} = A s^{(i)}, \quad x^{(i)} \in \mathbb{R}^w$$

$x_j^{(i)}$: recording of micro j at time i
($j=1 \dots n$)

$$x_j^{(i)} = \sum_k a_{jk} s_k^{(i)}$$

find $w = A^{-1}$ so that $s^{(i)} = w \cdot x^{(i)}$

$$w x = \begin{bmatrix} -w_1^T \\ \dots \\ -w_n^T \end{bmatrix} \cdot x$$

Lecture 14: Expectation - Maximization Algorithms

+) Clustering (k -means)

Data: $\langle x^{(0)}, \dots, x^{(m)} \rangle$

1. Initialize cluster centroids

$\mu_1, \dots, \mu_k \in \mathbb{R}^n$ randomly

2. Repeat until convergence:

(a) Set $c^{(i)} = \arg \min_j \|x^{(i)} - \mu_j\|_2^2$ ("color the points")

(b) for $j = 1 \dots k$

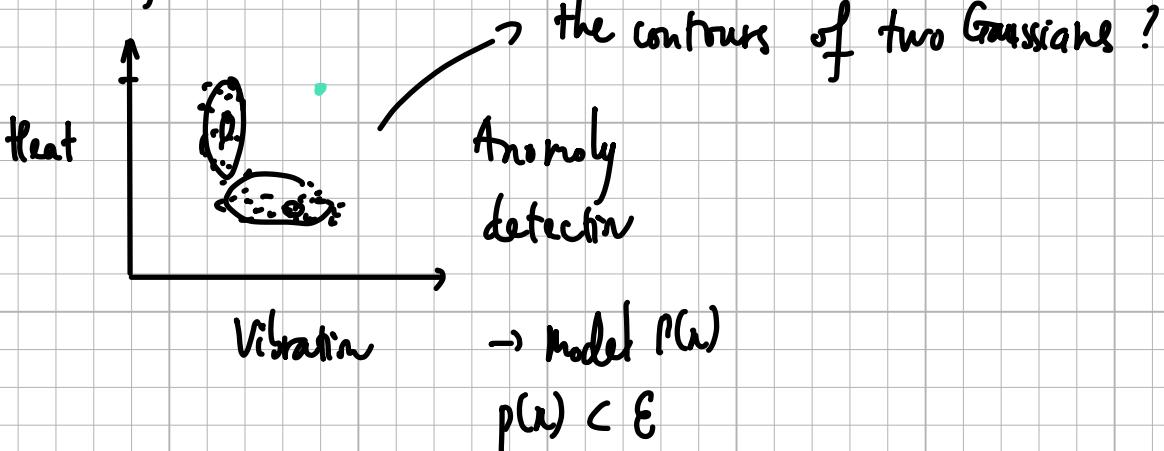
$$\mu_j := \frac{\sum_{i=1}^m \mathbf{1}\{c^{(i)} = j\} x^{(i)}}{\sum_{i=1}^m \mathbf{1}\{c^{(i)} = j\}} \quad (\text{"move the cluster to"})$$

$$J(c, \mu) = \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

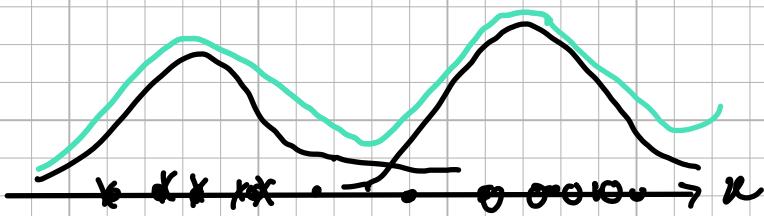
← assignments
Centroid

\Rightarrow purpose to choose k ?

+) Density estimation



+) 1-1 example ($\kappa \in \mathbb{R}$)



→ Gaussian discriminant analysis

→ EM algo !!! (expectation maximization)

+) Mixture of Gaussian model

Suppose there's a latent (hidden / unobserved)

random variable z , and $x^{(i)}, z^{(i)}$ are distributed

$$p(x^{(i)}, z^{(i)}) = p(x^{(i)} | z^{(i)}). p(z^{(i)})$$

where $z^{(i)} \sim \text{Multinomial}(\phi)$

$$x^{(i)} | z^{(i)} = j \sim N(\mu_j, \Sigma_j)$$

- If we knew the $z^{(i)}$, can use MLE

$$\ell(\phi, \mu, \Sigma) = \sum_{i=1}^n \log p(x^{(i)}, z^{(i)}; \phi, \mu, \Sigma)$$

$$\hat{\phi}_j = \frac{1}{n} \sum_{i=1}^n \mathbf{1}\{z^{(i)} = j\}$$

$$\hat{\mu}_j = \frac{\sum_{i=1}^n \mathbf{1}\{z^{(i)} = j\} x^{(i)}}{\sum_{i=1}^n \mathbf{1}\{z^{(i)} = j\}}$$

+) E-step (Guess value of $z^{(i)}$'s)

$$\text{Set } w_j^{(i)} = p(z^{(i)} = j | x^{(i)}; \phi, \mu, \Sigma)$$

$$= \frac{P(x^{(i)} | z^{(i)} = j) P(z^{(i)} = j)}{\sum_{l=1}^K P(x^{(i)} | z^{(i)} = l) P(z^{(i)} = l)} \sim \phi_j$$

$\mathcal{N}(\mu_j, \sigma_j)$

+ M-Step

$$\phi_j := \frac{1}{m} \sum_{i=1}^m w_j^{(i)}$$

ϕ_j : lecture notes

$$\mu_j = \frac{\sum_{i=1}^m w_j^{(i)} x^{(i)}}{\sum_{i=1}^m w_j^{(i)}}$$

$$E[1\{z^{(i)} = j\}]$$

$$P(x, z; \theta) \rightarrow \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^m P(z^{(i)}; \theta)$$

- $w_j^{(i)}$ is how much $x^{(i)}$ is assigned to the μ_j Gaussian

+ Jensen's Inequality

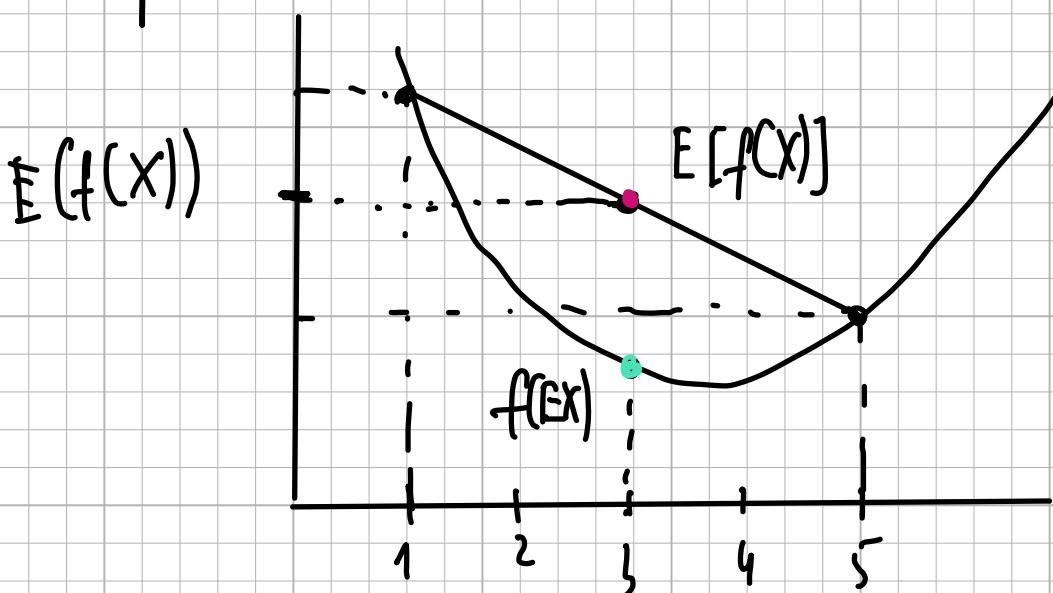
let f be a convex function (e.g. $f''(x) > 0$)

X be a random variable

$$\text{Then } f(E[X]) \leq E[f(X)]$$

$$X = \begin{cases} 1 & \text{with prob 1/2} \\ 3 & \text{with prob 1/2} \end{cases} \quad E[X] = \sum x_i p_i = 1 \cdot 0.5 + 3 \cdot 0.5 = 2$$

Example:



Điều " = " xảy ra khi f là strictly convex

$\Leftrightarrow X$ is constant

($X = \text{EX}$ with prob 1)

- Concave function

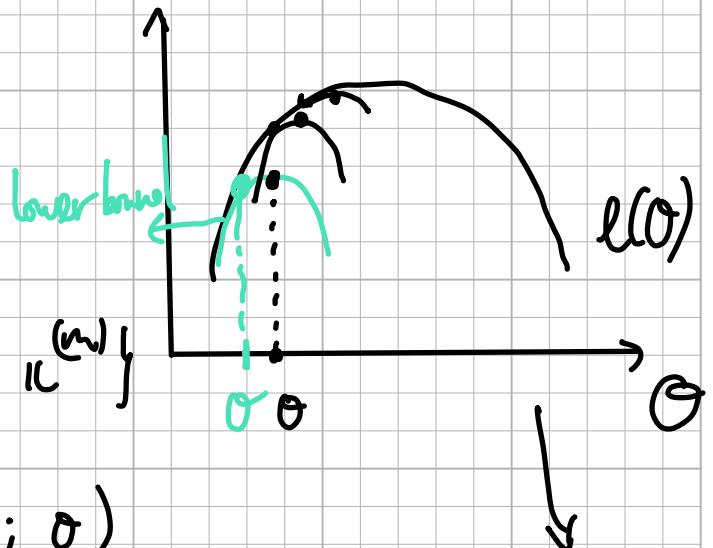
Have model for

$$P(x, z; \theta)$$

only to serve $x \in \{x^{(1)}, \dots, x^{(m)}\}$

$$l(\theta) = \sum_{i=1}^m \log P(x^{(i)}; \theta)$$

$$= \sum_{i=1}^m \log \sum_{z^{(i)}} P(x^{(i)}, z^{(i)}; \theta)$$



Converge
to local
optimum

Want $\underset{\theta}{\operatorname{argmax}} l(\theta)$

$$\underset{\theta}{\operatorname{Max}} \sum_i \log P(x^{(i)}; \theta)$$

$$= \sum_i \log \sum_{z^{(i)}} P(x^{(i)}, z^{(i)}; \theta)$$

$$= \sum_i \log \sum_{z^{(i)}} Q_i(z^{(i)}) \frac{P(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})}$$

where $Q(z^{(i)})$ is a prob distribution $\sum_{z^{(i)}} Q(z^{(i)}) = 1$

$$= \sum_i \log \mathbb{E}_{z^{(i)} \sim Q_i} \left[\frac{P(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})} \right]$$

↑ how to come up?

$$\geq \sum_i \mathbb{E}_{z^{(i)} \sim Q_i} \left[\log \frac{P(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})} \right]$$

$$= \sum_i \sum_{z^{(i)}} Q_i(z^{(i)}) \cdot \log \frac{P(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})}$$

On a given iteration of EM with parameters θ ,

we want : $Q_i(z^{(i)}) = ?$ xra $k^{(i)}$

$$\frac{P(x^{(i)}, z^{(i)})}{Q_i(z^{(i)})} = \text{constant}$$

Set $Q_i(z^{(i)}) \propto p(x^{(i)}, z^{(i)}; \theta)$

$$Q_i(z^{(i)}) = \frac{p(x^{(i)}, z^{(i)}; \theta)}{\sum_{z_j^{(i)}} p(x^{(i)}, z_j^{(i)}; \theta)} = \dots = p(z^{(i)} | x^{(i)}; \theta)$$

- E-step:

$$\text{Set } Q_i(z^{(i)}) := p(z^{(i)} | x^{(i)}; \theta)$$

- M-step:

$$\theta := \operatorname{argmax}_{\theta} \sum_i \sum_{z^{(i)}} Q_i(z^{(i)}) \log \frac{p(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})}$$

Lecture 15. EM Algorithm & Factor Analysis

Recap: $p(x^{(i)}, z^{(i)}) = p(x^{(i)} | z^{(i)}) \cdot p(z^{(i)})$

$\sum_{z^{(i)}}$ ~ Multinomial (ϕ)

$$(p(z^{(i)} = j) = \phi_j)$$

$$x^{(i)} | z^{(i)} = j \sim \mathcal{N}(\mu_j, \Sigma_j)$$

\nwarrow latent

E-step:

$$w_j^{(i)} = Q_i(z^{(i)} = j) = p(z^{(i)} = j | x^{(i)}; \phi; \mu; \Sigma)$$

" $p(z^{(i)} = j)$ "

M-step:

$$\underset{\phi, \mu, \Sigma}{\text{Max}} \sum_i \sum_{z^{(i)}} Q_i(z^{(i)}) \log \frac{P(x^{(i)}, z^{(i)}; \phi, \mu, \Sigma)}{Q_i(z^{(i)})}$$

$$= \sum_i \sum_j w_j^{(i)} \log \frac{1}{(2\pi)^{D/2} |\Sigma_j|^{1/2}} \exp \left(-\frac{1}{2} (x^{(i)} - \mu_j) \Sigma_j^{-1} (x^{(i)} - \mu_j) \right) \cdot \phi_j$$

$$\nabla_{\mu_j} (\dots) \text{ set } 0$$

$$\Rightarrow \mu_j = \frac{\sum_i w_j^{(i)} x^{(i)}}{\sum_i w_j^{(i)}}$$

$w_j^{(i)}$ "strength" with which
 $x^{(i)}$ is assigned
to Gaussian j

$$\nabla_{\phi} (\dots) \text{ set } 0, \quad \nabla_{\Sigma} (\dots) \text{ set } 0$$

optimal
value

$$\phi_j = \frac{\sum_i w_j^{(i)}}{\sum_i w_j^{(i)}}$$

$$+) \quad \mathcal{T}(\theta, Q) = \sum_i \sum_{z^{(i)}} Q_i(z^{(i)}) \cdot \log \frac{P(x^{(i)}, z^{(i)}; \theta)}{Q_i(z^{(i)})}$$

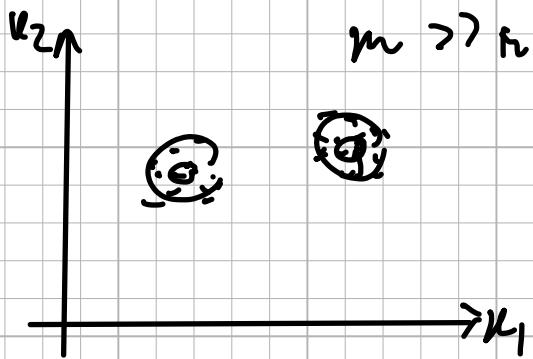
we know $\lambda(\theta) \geq \mathcal{T}(\theta, Q)$ for any θ, Q

E-step: Maximize \mathcal{T} wrt Q

M-step: Maximize \mathcal{T} wrt θ

+) Mixture of Gaussian

Say $n=2$, $m=100$



$m \approx n$, or $m \ll n$

Say $m=10$, $n=100$

\downarrow \downarrow
 50 days 100 sensors

$$\mu = \begin{bmatrix} \cdot \\ \cdot \\ \vdots \\ \cdot \end{bmatrix}_{100}^T \rightarrow \text{anomaly}$$

$P(x) \leq \epsilon$

Model as single Gaussian

$$x \sim N(\mu, \Sigma)$$

$$\text{MLE: } \mu = \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

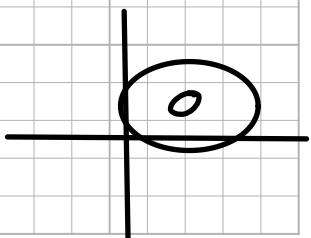
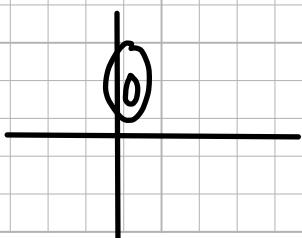
Maximum likelihood estimate

$$\Sigma = \frac{1}{m} \sum_{i=1}^m (x^{(i)} - \mu)(x^{(i)} - \mu)^T$$

If $m \leq n \Rightarrow \Sigma$ is singular / non invertible

Option 1: Constrain Σ to be diagonal

$$\Sigma = \begin{bmatrix} \sigma_1^2 & 0 & \dots & 0 \\ 0 & \sigma_2^2 & \dots & \vdots \\ \vdots & \ddots & \ddots & \ddots \\ 0 & \dots & \dots & \sigma_n^2 \end{bmatrix} \rightarrow \text{axes-aligned contours}$$



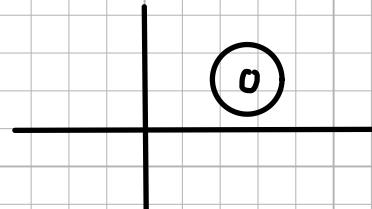
NLE: $\hat{\sigma}_j^2 = \frac{1}{m} \sum_i (x_j^{(i)} - \mu_j)^2 \Rightarrow$ Assume uncorrelated
n parameters between any 2s

Option 2

$$\text{Constrain } \Sigma = \begin{bmatrix} \sigma^2 & \dots & \dots \\ \dots & \ddots & \dots \\ \dots & \dots & \sigma^2 \end{bmatrix}$$

1 parameters \Rightarrow circular contours

$$\text{NLE: } \hat{\sigma}^2 = \frac{1}{m} \frac{1}{n} \sum_i \sum_j (x_j^{(i)} - \mu_j)^2$$



$$+) P(x_i | z) = P(x_i | z) \cdot P(z)$$

z is hidden

$$z \sim \mathcal{N}(0, 1) \quad d=5$$

$$n=ln \quad h=50$$

Factor analysis model

$$z \in \mathbb{R}^d \quad (d \times n)$$

$$x = \mu + \lambda z + \epsilon \quad \xrightarrow{\text{"pri"}} \text{Equivalently}$$

$$\text{where } \epsilon \sim \mathcal{N}(0, \Sigma) \quad x/z \sim \mathcal{N}(\mu + \lambda z, \epsilon)$$

- Parameters

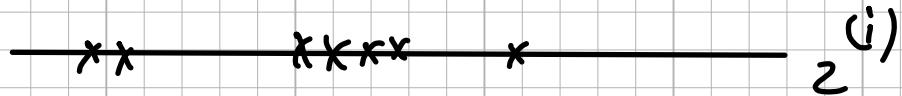
$$\mu \in \mathbb{R}^n, \lambda \in \mathbb{R}^{n \times d}, \Sigma \in \mathbb{R}^{n \times n \text{ diagonal}}$$

\hookrightarrow noise between seasons
independent

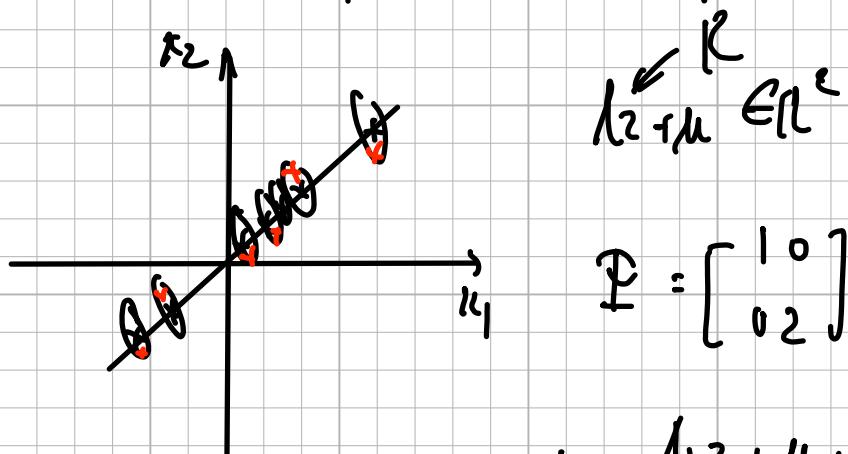
Examples

1. $z \in \mathbb{R}^1$, $x \in \mathbb{R}^2$

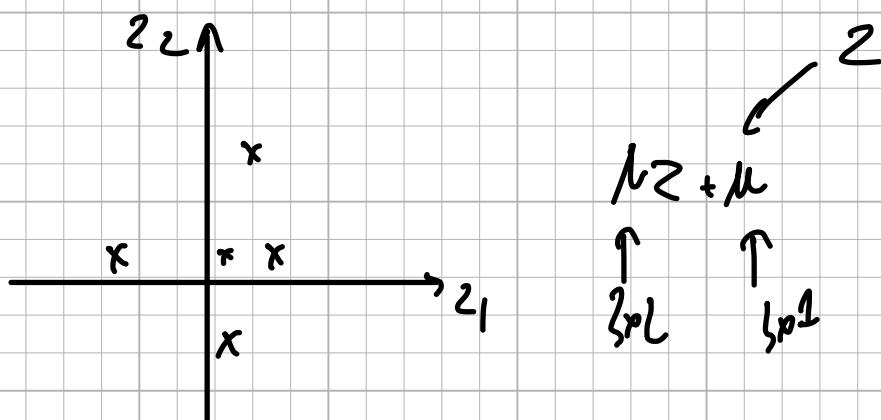
$d=1 \rightarrow$ lies on 1D subspace with mizes
 $n=2 \rightarrow$ 2D data
 $m=7$ $z \sim \mathcal{N}(0, 1)$



say $\Lambda = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$, $\mu = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$



2. $z \in \mathbb{R}^2$, $x \in \mathbb{R}^1$, $m=5$, $d=1$, $n=3$



+) Multivariate Gaussian

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \stackrel{\text{Def}}{\sim} \mathcal{N}(\mu, \Sigma)$$

$$x_1 \in \mathbb{R}^r, x_2 \in \mathbb{R}^s$$

$$x \in \mathbb{R}^{r+s}$$

$$x \sim N(\mu, \Sigma)$$

$$\mu = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \quad \begin{matrix} \uparrow r \\ \downarrow s \end{matrix}$$

Marginal $P(x_i) = ?$

$$P(x) = P(x_1, x_2)$$

$$= \int_{x_2} P(x_1, x_2) dx_2 \cdot P(x_1)$$

$$P(x_1) \text{ given } x_1 \sim \mu(u_1, \Sigma_{11})$$

$$\text{Conditional } P(x_1 | x_2) = ?$$

$$x_1 | x_2 \sim \mathcal{N}(\mu_{1|2}, \Sigma_{1|2})$$

$$\frac{P(x_1, x_2)}{P(x_2)}$$

$$\mu_{1|2} = \mu_1 + \Sigma_{12} \Sigma_{22}^{-1} (x_2 - \mu_2)$$

$$\Sigma_{12} = \Sigma_{11} - \Sigma_{12} \Sigma_{22}^{-1} \Sigma_{21}$$

Step 1: Derive $P(x_1, z)$

$$\binom{2}{\mu} \sim \mathcal{N}(\mu_{1|2}, \Sigma)$$

$$z \sim \mathcal{N}(0, 1)$$

$$Ez = 0$$

$$x = \mu + \lambda z + \epsilon$$

$$Ex = E(\mu + \lambda z + \epsilon)$$

$$= \mu$$

$$\mu_{x2} = \begin{bmatrix} 0 \\ \mu \end{bmatrix} \begin{array}{c} \uparrow d \\ \uparrow n \end{array}$$

$$\bar{\Sigma}_{22} = E(x - \bar{x}_2)(x - \bar{x}_2)^T$$

$$\bar{\Sigma} = \begin{bmatrix} \bar{\Sigma}_{11} & \bar{\Sigma}_{12} \\ \bar{\Sigma}_{21} & \bar{\Sigma}_{22} \end{bmatrix} \begin{array}{c} \hookrightarrow \\ \downarrow d \end{array} \begin{array}{c} \hookrightarrow \\ \downarrow n \end{array}$$

$$= E(\Lambda_{22}^T \Lambda^T) + E(\epsilon \epsilon^T)$$

$$= N \bar{\Sigma}_{22}^T \bar{\Sigma}^T + \bar{\Sigma}$$

$$= M^T + \bar{\Sigma}$$

$$:= \begin{bmatrix} I & \Lambda^T \\ \Lambda & \Lambda \Lambda^T + \Sigma \end{bmatrix}$$

$$\begin{bmatrix} z \\ \epsilon \end{bmatrix} \sim \mathcal{N} \left(\begin{bmatrix} 0 \\ \mu \end{bmatrix}, \begin{bmatrix} I & \Lambda^T \\ \Lambda \Lambda^T + \Sigma \end{bmatrix} \right)$$

$$P(\mu^{(i)})$$

• E-Step:

$$Q_i(z^{(i)}) = P(z^{(i)} | x^{(i)}; \theta)$$

$$z^{(i)} | x^{(i)} \sim \mathcal{N}(\mu_{z^{(i)} | x^{(i)}}, \Sigma_{z^{(i)} | x^{(i)}})$$

where

$$\mu_{z^{(i)} | x^{(i)}} = \vec{0} + \Lambda^T (\Lambda \Lambda^T + \Sigma)^{-1} (x^{(i)} - \mu)$$

$$\Sigma_{z^{(i)} | x^{(i)}} = I - \Lambda^T (\Lambda \Lambda^T + \Sigma)^{-1} \Lambda$$

• M-step:

$$Q_i(z^{(i)})$$

$$\theta := \arg \max_{\Theta} \sum_i \int_{z^{(i)}} Q_i(z^{(i)}). \log \frac{P(x^{(i)}, z^{(i)})}{Q_i(z^{(i)})} dz^{(i)}$$

$$= \sum_i \mathbb{E}_{z^{(i)} \sim Q_i} \left[\underbrace{\log \frac{P(x^{(i)}, z^{(i)})}{Q_i(z^{(i)})}}_{\text{quadratic function}} \right]$$

$$\nabla_{\mu} = 0$$

Lecture 6: Independent Component Analysis & RL

ICA Problem:

Sources: $s \in \mathbb{R}^n$

$s_j^{(i)}$ = speaker j at time i

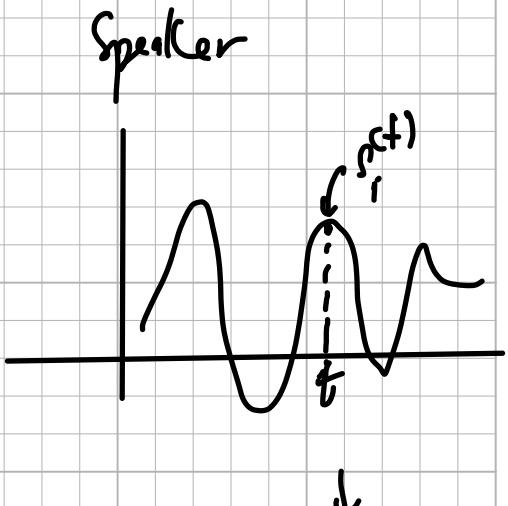
$$x^{(i)} = A s^{(i)}, x \in \mathbb{R}^r$$

Goal: Find $W = A^{-1}$, so

$$s^{(i)} = W x^{(i)}$$

$$W = \begin{bmatrix} -W_1^T - \\ \vdots \\ -W_n^T - \end{bmatrix}$$

"non Gaussian"



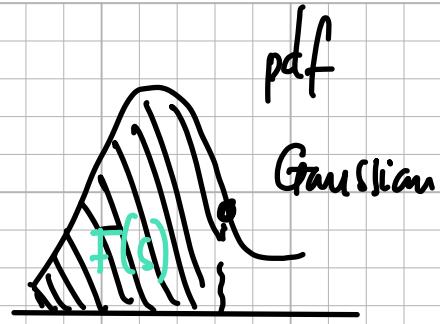
$$p_s(s)$$

Cumulative DF: constant

$$F(s) = P(S \leq s)$$

random variable

e.g.



$$p_s(s) = F'(s)$$

b.t.c.: Choose $F(s)$

$$F(s) = 1 \{ 0 \leq s \leq 1 \} (s \sim \text{Uniform}(0,1))$$

Density of s : $p_s(s)$

$$\text{say } x = 2s \quad (A=2, w=\frac{1}{2})$$

$$x = As = w^{-1}s$$

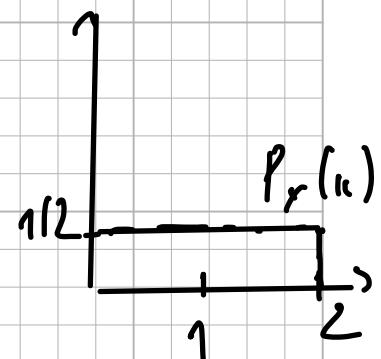
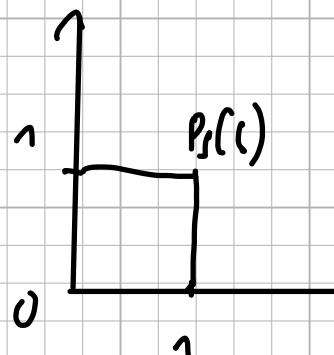
$$x \sim \text{Uniform}(0,2)$$

$$s = wx$$

Pertinent of w

$$p_x(x) = p_s(wx) |w|$$

Density of x



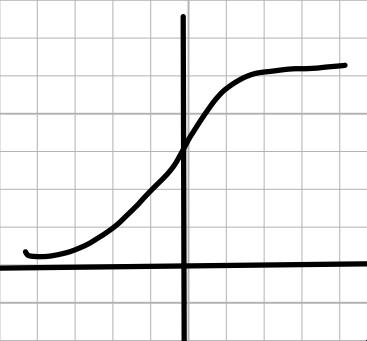
$$p_x(x) = \frac{1}{2} \mathbb{1}_{\{0 \leq x \leq 2\}}$$

$$p_s(s) = ?$$

$$F(s) = P(S \leq s)$$

↑ chance of person voice

→ Laplace distribution

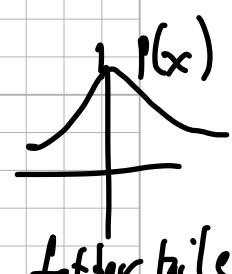


also works

sigmoid function

$$F(s) = \frac{1}{1 + e^{-s}}$$

$$p_s(s) = F'(s)$$



fatter tails

$$P_S(s) = \prod_{i=1}^n P_{S_i}(s_i)$$

(n speakers independent)

$$\begin{aligned} p(x) &= P_S(w_x) |w| \\ &= \left(\prod_{j=1}^m P_{S_j}(w_j^T x) \right) |w| \end{aligned}$$

$$g(z) = \frac{1}{1+z^{-2}}$$

MLE:

$$\ell(w) = \sum_{i=1}^n \left[\ln \left(\prod_j P_S(w_j^T x^{(i)}) \right) \right] / |w|$$

Stochastic gradient descent:

$$\nabla_w \ell(w) = \begin{bmatrix} 1 - g(w_1^T x) \\ \vdots \\ 1 - g(w_n^T x) \end{bmatrix} w^T (w^T)^{-1}$$

\rightarrow Recover $s = wx$

+) Reinforcement Learning
 \rightarrow Cost function, Reward function

$$\begin{array}{ll} R(s) = +1 & \text{for win} \\ \uparrow & \uparrow \\ \text{Reward} & \text{State} \end{array} \quad \begin{array}{ll} -1 & \text{for lose} \\ 0 & \text{for tie} \end{array}$$

Credit assignment problem

+) Markov Decision Process (MDP)

$$(S, A, \{P\}, \gamma, R)$$

S = set of states

A = set of actions

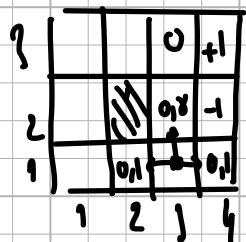
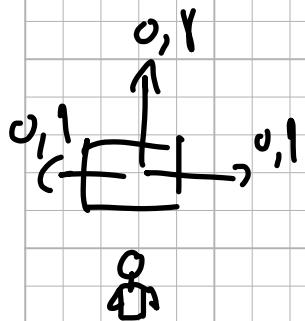
P_{SA} - state transition probabilities γ - discount factor

$$\left(\sum_{S_i} P_{SA}(s_i) = 1 \right)$$

$$\gamma \in [0, 1]$$

R - reward function

- MDP - Robot navigating this simple maze, a obstacle



Actions {N, S, E, W}

$$P_{(3,1)N}((3,2)) = 0.8$$

$$R((4,1)) = +1$$

$$0,1 = P_{(2,1)}((3,1)) = P_{(3,1)}((4,1))$$

$$R((4,1)) = -1$$

$$P_{(1,1)}((3,1)) = 0$$

$P(S) = -0.02$ for all other states. Total payoff

So

$$R(S_t) + \gamma R(S_1) + \gamma^2 R(S_2) + \dots$$

Choose action a_0

$$\gamma = 0.95$$

Get to $S_1 \sim P_{S_0 a_0}$

Choose action a_1

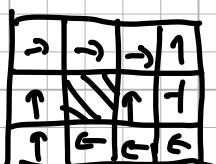
Get to $S_2 \sim P_{S_1 a_1}$

:

Policy $\pi: S \rightarrow A$
(Controller)

when in state 3,
take action $\pi(S)$

optimal policy:



Lecture 17: MDP & Value / Policy Iteration

i) Define: $\underline{V}^n, V^*, \pi^*$

for a policy π , $V^n : S \rightarrow \mathbb{R}$ s.t.

s.t. $V^n(s)$ is the expected total payoff
from starting in state s and executing π

$$V^n(s) = E \left[R(S_0) + \gamma R(S_1) + \dots \mid \pi, S_0 = s \right]$$

Immediate reward

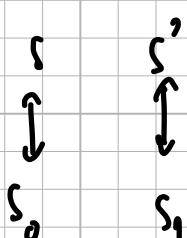
"Value function for policy π "

		π		
		?		
?	→	→	→	+1
2	↓	↖ ↘	→	-1
1	→	→	↑	↑
	1	2	3	4

		V^n
		?
.	.52	.75
.	-.90	.77
.	-.88	-1
.	-.87	-1.20

ii) Bellman's Equations:

$$V^n(s) = R(s) + \gamma \sum_{s'} P_{S^n(s)}(s') V^n(s')$$



$$V^n(s) = E \left[R(s) + \gamma V^n(s') \mid \begin{matrix} s \\ s' \end{matrix} \right]$$

$$s' \sim P_{S^n(s)}$$

In state s , take action $a : \pi(s)$, $s' \sim P_{S^a}$

Given n , got a linear system of equations in terms of $V^n(s)$

$$V^n((1,1)) = R((1,1)) + \gamma (0.8 V^n((3,1)) + 0.1 V^n((2,1)) + 0.1 V^n((4,1)))$$

- 11 states $\rightarrow 11$ vectors, linear solver

+) V^* is the optimal value function

$$V^*(s) = \max_n V^n(s)$$

Bellman equation: $V^*(s) = R(s) + \max_a \gamma \sum_{s'} P_{sa}(s') V^*(s')$

expected future reward

$$\pi^*(a) = \arg \max_a \sum_{s'} P_{sa}(s') V^*(s')$$

optimal policy

+) Practice w/ confusing notation:

$$V^*(s) = V^{n''}(s) \geq V^n(s)$$

for every n, s

Strategy:
 1) find V^*
 2) use argmax func to find π^*

+) Value Iteration:

Initialize $V(s) := 0$ for every s

For every s , update

$$V(s) := R(s) + \max_a \gamma \sum_{s'} P_{sa}(s') V(s')$$

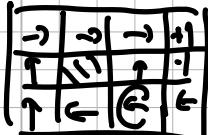
"synchronous update" vs. "asynchronous update"

Bellman backup operator

$$V = B(V)$$

Value Iteration cause V to converge to V^*

3	.86	.90	.95	-1
2	.81	.81	.65	-1
1	.78	.75	.71	.41



$$\begin{aligned} "E" &= 0,8 \cdot 0,75 + 0,1 \cdot 0,69 + 0,1 \cdot 0,49 \\ &= 0,740 \end{aligned}$$

$$\begin{aligned} "P" &= 0,8 \cdot 0,69 + 0,1 \cdot 0,75 + 0,1 \cdot 0,49 \\ &= 0,690 \end{aligned}$$

+) Policy iteration

Initialize π randomly

Repeat:

Set $V := V^n$ (i.e., solve Bellman's update to get V^n)

Set $\pi(s) := \arg \max_a \sum_{s'} P_{sa}(s') V(s')$

V^n - linear system of equations

↓
small problems: faster
expensive

+) What if we don't know P_{sa} :

$P_{sa}(s') = \frac{\# \text{times took action } "a" \text{ in stat } s \text{ and got to } s'}{\# \text{times took action } "a" \text{ in state } s}$

$\# \text{times took action } "a" \text{ in state } s$

(or $\frac{1}{|S|}$ if above is $\frac{0}{0}$) $\xrightarrow{\text{if}} \frac{1}{|A|}$ (replace)

+1 Putting it together:

Repeat: {

ϵ -greedy

\uparrow 0,1 chance wrt π
0,1 chance randomly

take actions wrt π to get experience on MDP

Update estimates of P_{SA} , (and possibly R)

Solve Bellman's equations using value iteration to get V

Update $\pi(s) = \arg \max_a \sum_{s'} P_{Sa}(s') V(s')$

}

+1 Exploration vs. Exploitation

"boltzmann exploration"

"intrinsic motivation"

DML

DL

Curiosity

Data mining

Lecture 18: Continuous State MDP & Model Simulation

- Discretization
- Models simulation
- fitted value iteration

+ Recap:

$$\text{MDP: } (S, A, \{P_{SA}\}, \{R\}, R)$$

$$V^n(s) = E[R(s_0) + \gamma R(s_1) + \dots | n, s=s_0]$$

$$V^*(s) = \max_n V^n(s)$$

$$n^*(s) = \arg \max_a \sum_{s'} P_{SA}(s') V^*(s')$$

$\hookrightarrow E_{s' \sim P_{SA}}[V^*(s')]$

Value iteration

$$V(s) := R(s) + \max_a \sum_{s'} P_{SA}(s') V(s')$$

roll, pitch, yaw

Car, / kinematic model
Helicopter } dynamic model

x, y, z

$x, y, z, \phi, \theta, \psi$

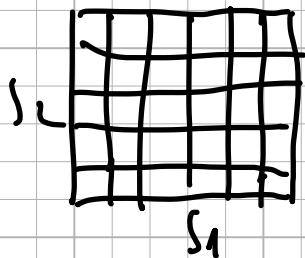
$\dot{x}, \dot{y}, \dot{z}$

$\dot{x}, \dot{y}, \dot{z}, \dot{\phi}, \dot{\theta}, \dot{\psi}$

free hinge : $x, \theta, \dot{x}, \dot{\theta}$ to keep balanced

State space: R^n

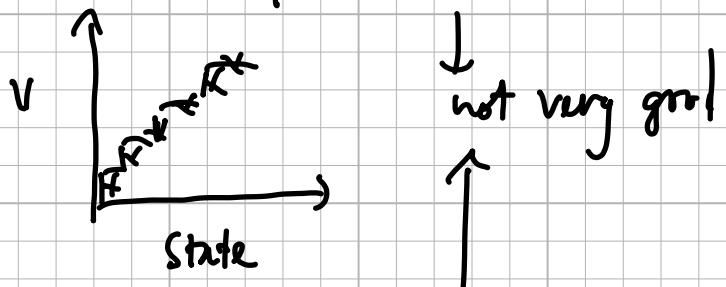
+) Discretization



1. P x_1, y : state inverted pendulum

Problems :

- Naive representation for V^* (and π^*)



- Curse of dimensionality

$S = \mathbb{R}^n$, and discretize

each dimensions into K values

get K^n discrete space

+) Approximate V^* directly, no resort to discretization

$$y \approx \theta^\top \phi(x), \quad \phi(x) : \text{feature of } x = \begin{bmatrix} x_1 \\ x_2 \\ x_1^2 \\ x_1 x_2 \\ \vdots \end{bmatrix}$$

$$V^*(s) \approx \theta^\top \phi(s)$$

+) Model (simulator) of MDP :

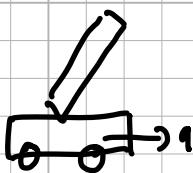
$$s \xrightarrow{\text{Model}} s' \sim p_{sa}$$

\uparrow
a (\approx discretize)

+) How to get model ?

- Physics simulator

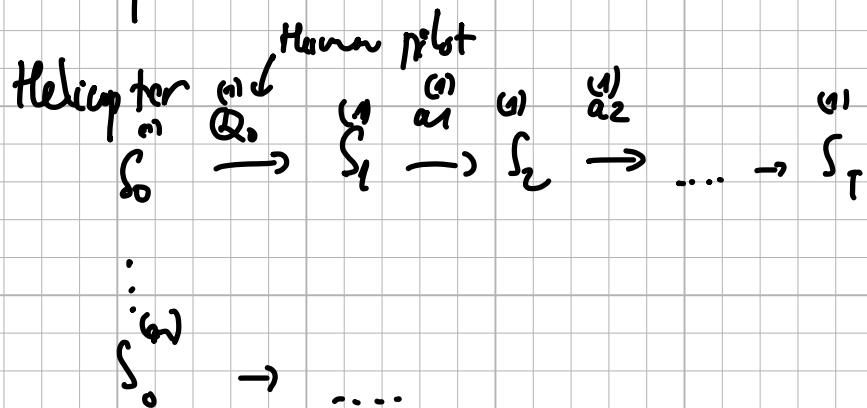
$$s' = s + (\Delta t) \cdot \dot{s}$$



$$\dot{s} = \begin{pmatrix} \dot{x} \\ \dot{\theta} \end{pmatrix}, \quad s = \begin{pmatrix} x \\ \theta \end{pmatrix}$$

$$\dot{s} = \begin{pmatrix} \dot{x} \\ \dot{\theta} \end{pmatrix} = \begin{pmatrix} \alpha - L_f \cos(\theta)/m \\ 0 \end{pmatrix}$$

- Learn from data



Apply supervised learning to estimate

s_{t+1} on function of s_t, a_t
(s, a)

e.g.: linear regression version (low-speed Helicopter)

$$s_{t+1} = A s_t + B a_t$$

\uparrow
 R^{new}

$$\min_{A, B} \sum_{i=1}^n \sum_{t=0}^T \| s_{t+1}^{(i)} - (A s_t^{(i)} + B a_t^{(i)}) \|^2$$

Model

$$s_{t+1} = A s_t + B a_t \quad (\text{deterministic})$$

or

$$s_{t+1} = A s_t + B a_t + \epsilon_t \quad \text{where } \epsilon_t \sim N(0, \sigma^2 I) \quad (\text{stochastic})$$

"Model-free RL"

"Model-based"

f) fitted value iteration:

Choose feature $\phi(s)$ of state s

$$V(s) = \Theta^T \phi(s)$$

$\phi(s)$ = non linear func...

+ Value Iteration

$$V(s) := R(s) + \gamma \max_a \sum_{s'} P_{s,a}(s') V(s')$$

$$= R(s) + \gamma \max_a \underbrace{E_{s' \sim P_{s,a}} [V(s')]}_{s \rightarrow V(s)}$$

fitted Value Iteration

Sample $\{s^{(1)}, \dots, s^{(n)}\} \subseteq S$ randomly

Initialize $\Theta := 0$

$$V(s) = \max_a \underbrace{E_{s \sim P_{s,a}} [R(s) + \gamma V(s')]}_{q(a)}$$

Repeat {

for $i = 1 \dots n$ {

For each action $a \in A$ \downarrow using model

Sample $s'_1, s'_2, \dots, s'_K \sim P_{s^{(i)}, a}$

$$\text{Set } q(a) := \frac{1}{K} \sum_{j=1}^K [R(s^{(i)}) + \gamma V(s'_j)]$$

}

Fitted VI:

$$\begin{aligned} \text{Want } V(s^{(i)}) &\approx y^{(i)} \\ \text{i.e. } \Theta^T \phi(s^{(i)}) &\approx y^{(i)} \end{aligned}$$

$$\text{Set } y^{(i)} = \max_a q(a) \quad \text{and} \quad \theta := \arg \min_{\theta} \frac{1}{2} \sum_{i=1}^m (\theta^T \phi(s^{(i)}) - y^{(i)})^2$$

- Fitted VI gives approximation to V^*

Implicitly defines π^*

$$\pi^*(s) = \arg \max_a E_{s' \sim P_{S|s}} [V^*(s')]$$

Used $s'_1, \dots, s'_L \sim P_{S|s}$ to approximate expectation

Say model is

$$s_{t+1} = f(s_t, a_t) + \epsilon_t$$

(e.g. $A s_t + B a_t + \epsilon_t$)

for deployment (running)

Set $\epsilon_t = 0$ and $K=1$

When on state s ,

Pick action \uparrow simulation w/o noise

$$\arg \max_a V(f(s, a))$$

\downarrow
 $s' \sim P_{S|s}$, with deterministic transition

Lecture 15: Reward Model & Linear Dynamical System

Outline : - State-action rewards

- finite horizon MDP

- linear dynamical systems
+ Model

+ LQR (Linear-quadratic regulation)

Recap: MDP $(S, A, \{P_{SA}\}, \gamma, R)$

Value iteration:

$$V(s) := R(s) + \max_a \gamma \sum_{s'} P_{Sa}(s') V(s')$$

$$\pi^*(s) = \arg \max_a \sum_{s'} P_{Sa}(s') V^*(s')$$

+ State-action rewards

$$R: S \times A \mapsto \mathbb{R}$$

$$R(S_0, a_0) + \gamma R(S_1, a_1) + \gamma^2 R(S_2, a_2) + \dots$$

Hellman's equals

$$V^*(s) = \max_a [R(s, a) + \gamma \sum_{s'} P_{Sa}(s') V^*(s')]$$

immediate reward future rewards

Value iteration $V(s) := \text{RHS}$

$$\pi^*(s) = \arg \max_a R(s, a) + \gamma \sum_s P_{s,a}(s') V(s')$$

+1 finite horizon MDP.

$$(S, A, \{P_{s,a}\}, T, R)$$

Horizon time T ($\text{Nb } 1$)

Payoff
 $E [R(S_0, a_0) + R(S_1, a_1) \dots + R(S_T, a_T)]$

$\pi_t^*(s)$: Non-stationary policy

↳ changes over time

$\pi^*(s)$: Stationary policy

Non-stationary State transitions

$S_{t+1} \sim P_{S_t, a_t}^{(t)}$ at certain state, certain time
 $R_{(s, a)}^{(t)}$

- Changing dynamics

- weather forecasts

- Industrial automation

...

Optimal Value function:

$$V_t^*(s) = E[R(s_t, a_t) + \dots + R(s_7, a_7) | n^*, s_0 = s]$$

expected total payoff starting in state s at time t , execute n^*

Value iteration (Dynamic Programming)

$$V_t^*(s) = \max_a R(s, a) + \sum_{s'} P_{s,a}^{(t)} V_{t+1}^*(s')$$

$$\pi_t^*(s) = \arg \max_a \text{_____}$$

non-stationary

$$V_1^*(s) = \max_a R(s, a) \quad (\text{after 1 ends})$$

$$\rightarrow V_1^*, V_2^*, \dots, V_0^*$$

+ Linear Quadratic Regulation (LQR)
only relatively small set

$$(S, A, \{P_{sa}\}, T, R)$$

$$S = \mathbb{R}^n, A = \mathbb{R}^{d \times d}$$

$$P_{sa} : S_{t+1} = A S_t + B a_t + w_t$$

matrix \downarrow \uparrow
 $\mathbb{R}^{n \times n}$ $\mathbb{R}^{n \times d}$

w is $w_t \sim \mathcal{N}(0, \tilde{\Sigma}_w)$
 less important

x_1, y_1, θ
 $\ddot{x}_1, \ddot{y}_1, \ddot{\theta}$

$$R(s, a) = - \left(S^T V s + a^T V a \right)$$

where $V \in \mathbb{R}^{m \times w}$, $V \in \mathbb{R}^{d \times d}$

$V, V \geq 0$ (positive-semi-definite)

implies $S^T V S \geq 0$, $a^T V a \geq 0$

Want $S \approx 0$

Choose $V = I$, $V = I$

$$R(s, a) = - (||S||^2 + ||a||^2)$$

Where to get A, B ?

$$S_0^{(0)} \xrightarrow{a_0^{(0)}} S_1^{(1)} \xrightarrow{\dots} S_T^{(G)}$$

$$S_0^{(m)} \xrightarrow{a_m^{(m)}}$$

$$S_{t+1} \approx A s_t + b a_t$$

$$\min_{A, B} \sum_{t=1}^m \sum_{t=0}^{t-1} \| S_{t+1} - (A s_t + B a_t) \|^2$$

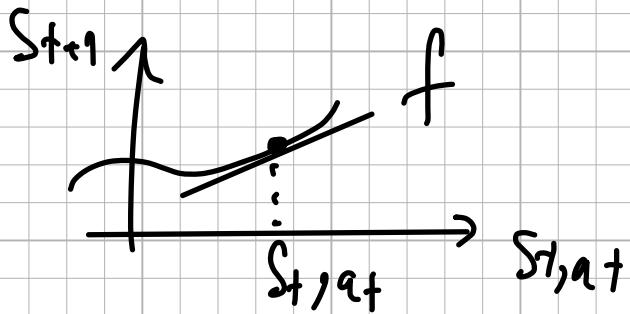
Method 1: learn from data

Method 2: linearize non-linear model

$$S_{t+1} = f(S_t, a_t)$$

$$\begin{pmatrix} x_{t+1} \\ \dot{x}_{t+1} \\ \theta_{t+1} \\ \dot{\theta}_{t+1} \end{pmatrix} = f \begin{pmatrix} x_t \\ \dot{x}_t \\ \theta_t \\ \dot{\theta}_t \end{pmatrix} + a_t$$

$$s_{t+1} = f(s_t, a_t)$$



$$s_{t+1} \approx \bar{s}_f + f'(\bar{s}_f)(s_t - \bar{s}_f)$$

$$\bar{s}_f \sim a \text{ const}$$

linearize f around (\bar{s}_f, \bar{a}_f) :

$$s_{t+1} \approx \bar{s}_f + f'(\bar{s}_f, \bar{a}_f) + \nabla_s f(\bar{s}_f, \bar{a}_f)^T (s_t - \bar{s}_f) + \nabla_a f(\bar{s}_f, \bar{a}_f)^T (a_t - \bar{a}_f)$$

$$\therefore s_{t+1} \approx A s_t + B a_t$$

$$s_t = \begin{pmatrix} 1 \\ x \\ \dot{x} \\ \theta \\ \dot{\theta} \end{pmatrix}$$

Problem:

$$S_{t+1} = A s_t + B a_t + \omega_t$$

$$R(s, a) = - (s^T U s + a^T V a)$$

Total payoff: $R(s_0, a_0) + \dots + R(s_T, a_T)$

+ Dynamic programming

$$V_T^*(s_T) = \max_{a_T} R(s_T, a_T)$$

$$= \max_{a_T} - s_T^T U s_T - a_T^T \underbrace{V_a}_{\geq 0}$$

$$= -s_T^T U s_T$$

$$V_T^*(s_T) = \overrightarrow{\sum} \quad R^{\text{new}}$$

Suppose

$$V_{t+1}^*(s_{t+1}) = s_{t+1}^T \Phi_{t+1}^{-1} s_{t+1} + \Psi_{t+1}^T$$

$$V_{t+1}^* \rightarrow V_t^*$$

$$V_t^*(s_t) = \max_{a_t} R(s_t, a_t) + E_{s_{t+1}} \sim p_{fa_t} \left[V_{t+1}^*(s_{t+1}) \right]$$

The best action is linear function | $N(A_{s+t}, B_{a_t}, \epsilon_t)$

$$V_t^*(s_t) = s_t^\top \vec{\Phi} s_t + \$_t$$

where Lecture Notes !!

+ To summarize:

Initialize $\Phi_T = -U$, $\Psi_T = 0$

Recursive calculate

Φ_t , Ψ_t using Φ_{t+1} , Ψ_{t+1} (for $t = T-1 \dots 0$)

Calculate L_t

$$\pi^*(s_t) = L_t s_t$$

V^* depends on Σ_w

π^* , L_t don't depend on Σ_w

Lecture 20: RL Debugging and Diagnostics

+ Debugging Improve model / simulator?

Diagnostics Modify reward func R?

Modify RL algorithm!

+) Policy search algorithm

("direct policy search")

$$V^* \rightarrow \pi^*$$

Try to find a good π directly

New definition:

A stochastic policy is a func

$\pi: S \times A \mapsto \mathbb{R}$ when $\pi(s, a)$ is the

prob of taking action a in state s

$$\sum_a \pi(s, a) = 1$$

→ Inverted pendulum

$$\pi_0(s, "R") = \frac{1}{1 + e^{-\theta^T s}}$$

$$\pi_0(s, "L") = 1 - \frac{1}{1 + e^{-\theta^T s}}$$

$$s = \begin{pmatrix} 1 \\ x \\ \dot{x} \\ \ddot{x} \\ \ddot{\dot{x}} \end{pmatrix}$$

Goal: find θ so that when we execute $\pi_0(s, a)$,
we minimize

$$\max_{\theta} E [R(s_0, a_0) + \dots + R(s_T, a_T) | \pi]$$

Not: Here, s_0 is a fixed initial state

+) Reinforce algorithm (let $T = 1$)

loop {

Sample $s_0, a_0, \dots, s_T, a_T$

Compute payoff = $R(s_0) + \dots + R(s_T)$

Update

$$\theta = \theta + \left[\frac{\pi_\theta(s_0, a_0)}{\pi_\theta(s_0, a_0)} + \dots \right] . \text{payoff}$$

}

+) POMDP
↳ partially observable

$$a = \theta^T s + \text{gaussian noise}$$

- at each step, get a partial (and potentially noisy) measurement at the state.

have to take action "a" using that

observation

$$y = \begin{pmatrix} x \\ f \end{pmatrix} + \text{noise}$$

$$V^*(s) \quad n^x(s)$$

+) low-level control task

$\int s^n$ or \sqrt{t} sampler